Gradient Descent and Backpropagation 1

Saeed Saremi

Assigned reading: 6.{2,4}, 7.{1,2}, 5.4.4

September 24, 2024

a summary of the previous lecture

• the geometry of θ in the logistic regression given the following parametrization: 1

$$p(1|x) = \frac{1}{1 + \exp(-\theta^{\top}x)} =: p(z(x;\theta)), \text{ where } z(x;\theta) = \theta^{\top}x$$

- θ is perpendicular to the decision boundary and points to class "1".
- ▶ softmax function: generalization of the logistic regression to multiclass (K > 2) classification via the generative approach, where we assumed $X|k \sim \mathcal{N}(\mu_k, \Sigma)$.
- ▶ the negative log-likelihood loss for the logisitic regression and its gradient:

$$\mathcal{L}(\theta) = -y \log p(z(x;\theta)) - (1-y) \log(1-p(z(x;\theta)))$$

$$\nabla_{\theta} \mathcal{L}(\theta) = (p_z - y) \nabla_{\theta} z = (p_z - y) x, \quad \begin{cases} \theta \leftarrow \theta - \xi \times \\ \theta \leftarrow \theta - \xi \times \end{cases} \quad \forall = 0$$

and we interpreted this gradient geometrically in terms of the parameter updates.

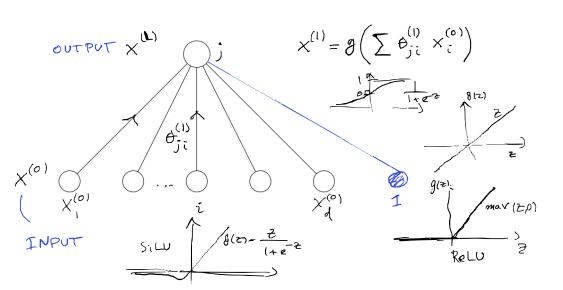
▶ artificial neural networks as simplified models of biological neural networks

¹Here, I assume the data is "centered", therefore $\theta_0 = 0$.

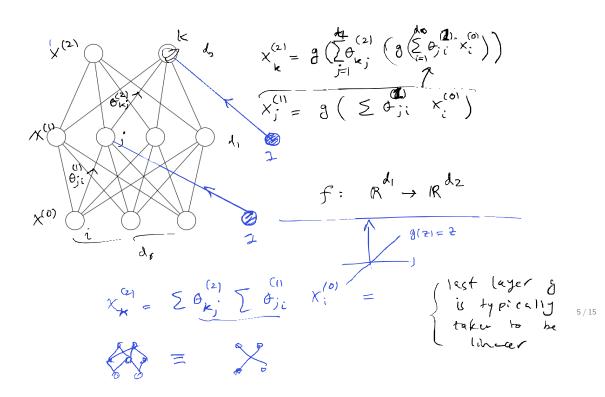
outline

- neural networks as a certain type of nonlinear functions
- > (linear/logisitic regression as single-layer neural network)
- > finish the comparisons with biological neural networks
- training neural networks: stochastic gradient descent
- > gradient of a function
- > the chain rule
- > convex and non-convex functions
- > geometrical interpretation of the gradient
- > stochastic gradient descent (SGD)
- **>** a prelude to backpropagation: the cross-entropy loss for multiclass (K > 2) classification and its gradient

single-layer neural networks L=1

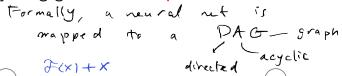


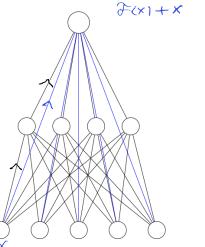
two-layer neural networks

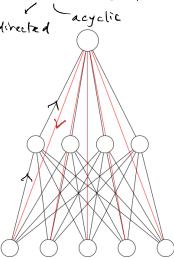


complex neural network architectures:

anything goes, but "loops" are not allowed!







biological neural networks v.s. deep neural networks



Figure: human brain: 10^{11} neurons, 10^{15} synapses to layers, feedback, (20 Hz, Spike timing

? What are we missing in this comparison?

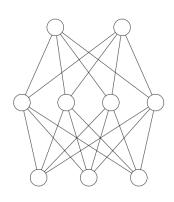
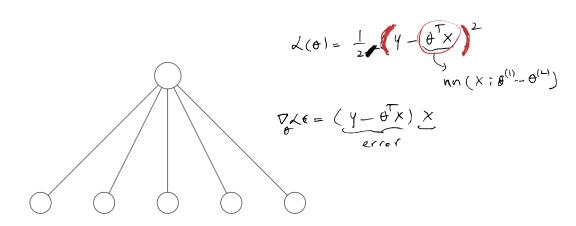
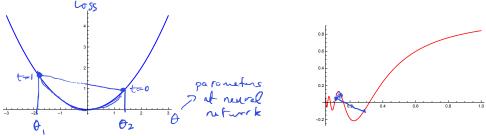


Figure: GPT-4: 10⁶ neurons, 10¹¹ parameters (weights), 100 layers PAG, GHZ, back prop

single-layer neural networks: the linear and logistic regression



convex and non-convex functions



Recall from the multivariate calculus that the gradient of the function $f: \mathbb{R}^m \to \mathbb{R}$ is the vector

$$\nabla f(\theta_1, \dots, \theta_m) = \left(\frac{\partial f}{\partial \theta_1}, \dots, \frac{\partial f}{\partial \theta_m}\right)^\top$$

- ▶ The gradient is the direction that leads to maximal increase of *f* (steepest ascent). Equivalently, the negative gradient is the direction of steepest descent.
- ▶ Formally, a function f is convex if for all θ_1 , θ_2 in \mathbb{R}^m and $t \in [0,1]$:

$$f(t\theta_1 + (1-t)\theta_2) \le f(\theta_1) + (1-t)f(\theta_2)$$

• convex functions have a single (global) minimum.

review: the chain rule from multivariate calculus

Given
$$u(z_1, z_2) = (z_1 + z_2)^2/2$$
, where $z_1(x_1, x_2) = x_1 \sin x_2$ and $z_2(x_1, x_2) = (\sin x_2)^2$; determine $\nabla_x u$ using the chain rule:

brute force:

orute force:

$$M(X_1, X_2) = \frac{1}{2} \left(X_1 \sin X_2 + \sin X_2 \right)^2 = \frac{1}{2} X_1^2 \sin X_2 + X_1 \sin X_2 + \frac{1}{2} \sin X_2$$

$$\frac{\partial M(X_{1,1} X_2)}{\partial X_2} = X_1^2 \sin X_1 \cos X_2 + \frac{3}{2} x_1 \sin X_2 \cos X_2 + \frac{2}{2} \sin X_2 \cos X_2$$

chain rule:
$$(z_1+z_2)$$
 (z_1+z_2)

$$\frac{\partial u}{\partial x_3} = \frac{\partial u}{\partial z_2} \frac{\partial z_1}{\partial x_2} + \frac{\partial u}{\partial z_2} \frac{\partial z_2}{\partial x_2}$$

$$= \left(z_1+z_2\right) \left(\frac{\partial z_1}{\partial x_2} + \frac{\partial z_2}{\partial x_2}\right)$$

$$= \left(\chi_1 \operatorname{Sm} x_2 + \operatorname{Fin} x_2\right) \left(\chi_1 \operatorname{cos} x_2 + 2 \operatorname{Sm} x_2 \operatorname{cos} x_2\right)$$

$$= \chi_1^2 \operatorname{Sm} x_2 \operatorname{coc} x_2 + 2 \chi_1 \operatorname{Sn} x_2 \operatorname{cos} x_2$$

$$+ \chi_1 \operatorname{Sm} x_3 \operatorname{coc} x_2 + 2 \operatorname{Sm} x_2 \operatorname{cos} x_2$$

$$+ \chi_2 \operatorname{Sm} x_3 \operatorname{coc} x_2 + 2 \operatorname{Sm} x_2 \operatorname{cos} x_2$$

the geometrical meaning of the gradient

Prove that $\nabla f(\theta)$, the gradient of the function $f:\Theta\to\mathbb{R}$ at any point θ , is perpendicular to the level set of f at that point.²

(Hint: define a path $t \to \theta$ on the level set and use the fact that by definition of the level set $\partial_t f(\theta(t)) = 0$. Use the chain rule!)

$$t \to \theta \to f$$

$$f(\theta(t)) = conit$$

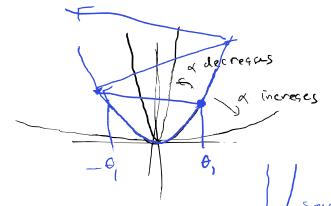
$$\partial_t f(\theta(t)) = 0$$

$$= \frac{\partial f}{\partial \theta_1} \frac{\partial \theta_1}{\partial t} + \cdots + \frac{\partial f}{\partial \theta_m} \frac{\partial \theta_m}{\partial t}$$

$$= (7f) \cdot \vec{V} = 0$$

²Recall that level sets are defined by $\{\theta': f(\theta') = f(\theta)\}$ for some constant c.

$$f(\theta) = \frac{\theta^2}{2\alpha}$$



$$\theta_{tH} = \theta_{t} - \epsilon \frac{\theta_{t}}{\alpha}$$

$$= (1 - \epsilon \frac{\theta_{t}}{\alpha}) \theta_{t}$$

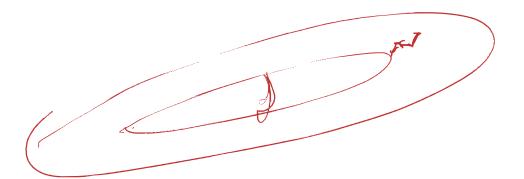
$$\theta_{t} = (1 - \epsilon \frac{\theta_{t}}{\alpha}) \theta_{t}$$



$$0 < E < \alpha =$$
 $(-\frac{E}{\alpha} < 1 =)$ As $+ +\infty = 0$

$$E = \alpha = \beta$$
 After only a single step
 $\theta_2 = \theta_{min}$

E> 2d



d_{mi} dictates the step size

stochastic gradient descent I

▶ at a high level the loss is always written as the sum of losses by individual points $i \in [n] := \{1, ..., n\}$ in the training set $\mathcal{D} = \{(x_i, y_i)\}_{i \in [n]}$:

$$\mathcal{L}(\theta) = \sum_{i=1}^{n} \mathcal{L}_{i}(\theta),$$
 $\nabla \mathcal{L}(\theta) = \sum_{i=1}^{n} \nabla \mathcal{L}_{i}(\theta)$

where $\mathcal{L}_i(\theta)$ is short for:

$$\mathcal{L}_i(\theta) := \mathcal{L}(x_i, y_i; \theta).$$

▶ This is very general, but it's very easy to see where it's coming from in the maximum log-likelihood framework. We always assume the i.i.d. setting:

$$(x_i, y_i) \stackrel{\text{iid}}{\sim} p_{\theta}(x, y), i \in [n].$$

Therefore,

$$p(\mathcal{D}|\theta) = \prod_{i=1}^{n} p_{\theta}(x_i, y_i).$$

It follows:

$$\mathcal{L}_i(\theta) = -\log p_{\theta}(x_i, y_i).$$

stochastic gradient descent II

The two prominent examples so far include:

▶ linear regression:

$$-\log p_{\theta}(x_i, y_i) = \frac{1}{2\sigma^2}(y_i - \theta^{\top} x_i^{(0)}) + const$$

▶ logistic "regression":

$$-\log p_{\theta}(x_{i}, y_{i}) = -y_{i} \log g(\theta^{\top} x_{i}^{(0)}) - (1 - y_{i}) \log(1 - g(\theta^{\top} x_{i}^{(0)})) + const$$

Note that the structure of the loss is general: we are getting ready to replace $\theta^{\top} x_i^{(0)}$ with an *L*-layer neural network: $f(x_i^{(0)}; \theta^{(1)}, \dots, \theta^{(L)})$.)

Stochastic Gradient Descent (SGD) in its pure form is defined by the following updates:

$$\theta_{t+1} = \theta_t - \epsilon_t \nabla_{\theta} \mathcal{L}(\mathbf{x}_i, \mathbf{y}_i; \theta),$$

where $i \in [n]$ is selected at random (typically without replacement) at each iteration t.

stochastic gradient descent III

- ▶ Intuitively (people have tried to study this) in high dimensions the loss landscape is "dominated" by saddle points and the noise in SGD helps to avoid them.
- ▶ SGD is convenient in the regime $n \gg 1$.
- One pass through the data is called an epoch.
- ▶ The problem is towards the end of the training (optimization) the noise in SGD will slow down the training.
- ▶ In short: noise helps us at the beginning of training, it "hurts" us towards the end.
- ▶ Of course, one can find a compromise by dividing the dataset into (random) mini-batches of size b: in this scheme one epoch involves $\lfloor n/b \rfloor$ updates.
- ? Based on this picture, can you suggest a batching scheme for effective training? ³

 $^{^3}$ Coming up with a mini-batch / learning rate schedule remains an art and problem dependent.